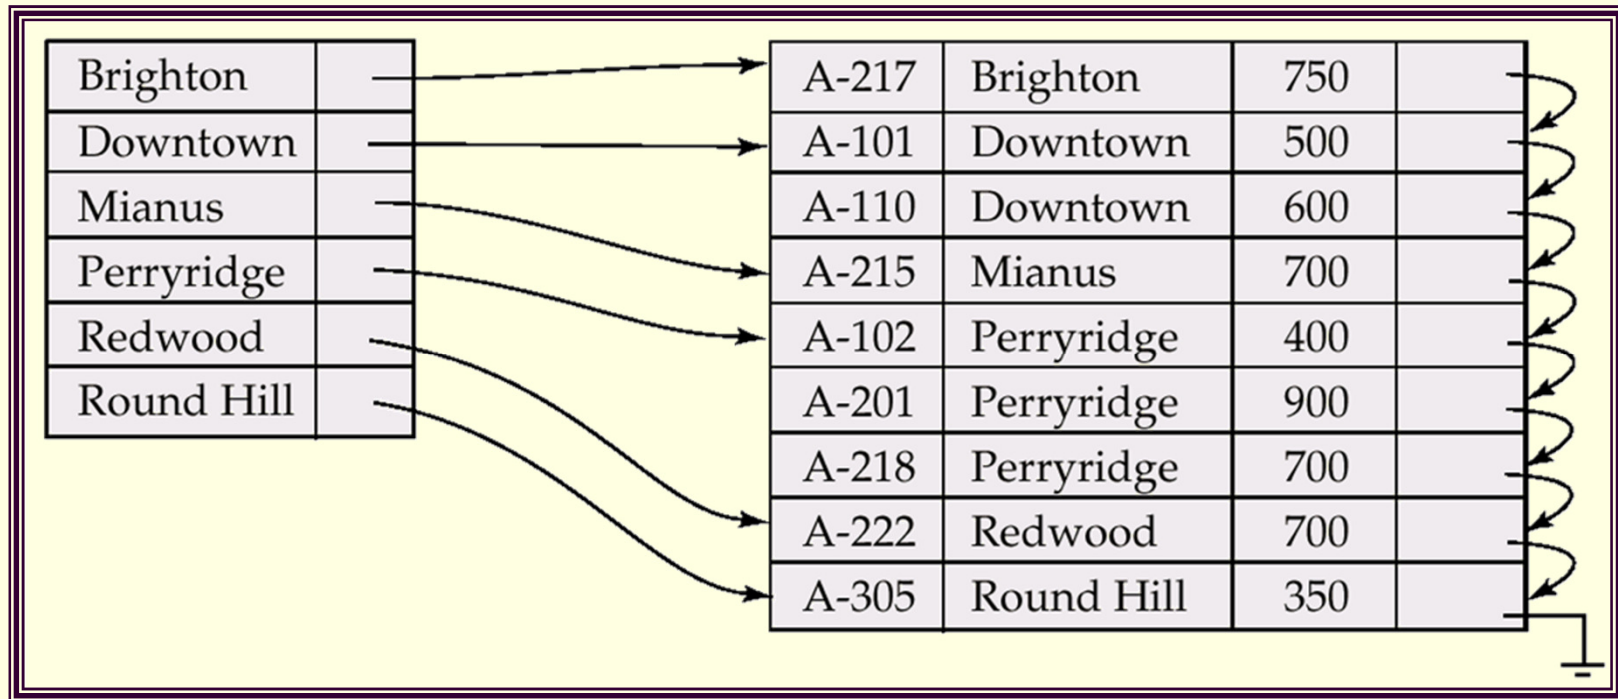


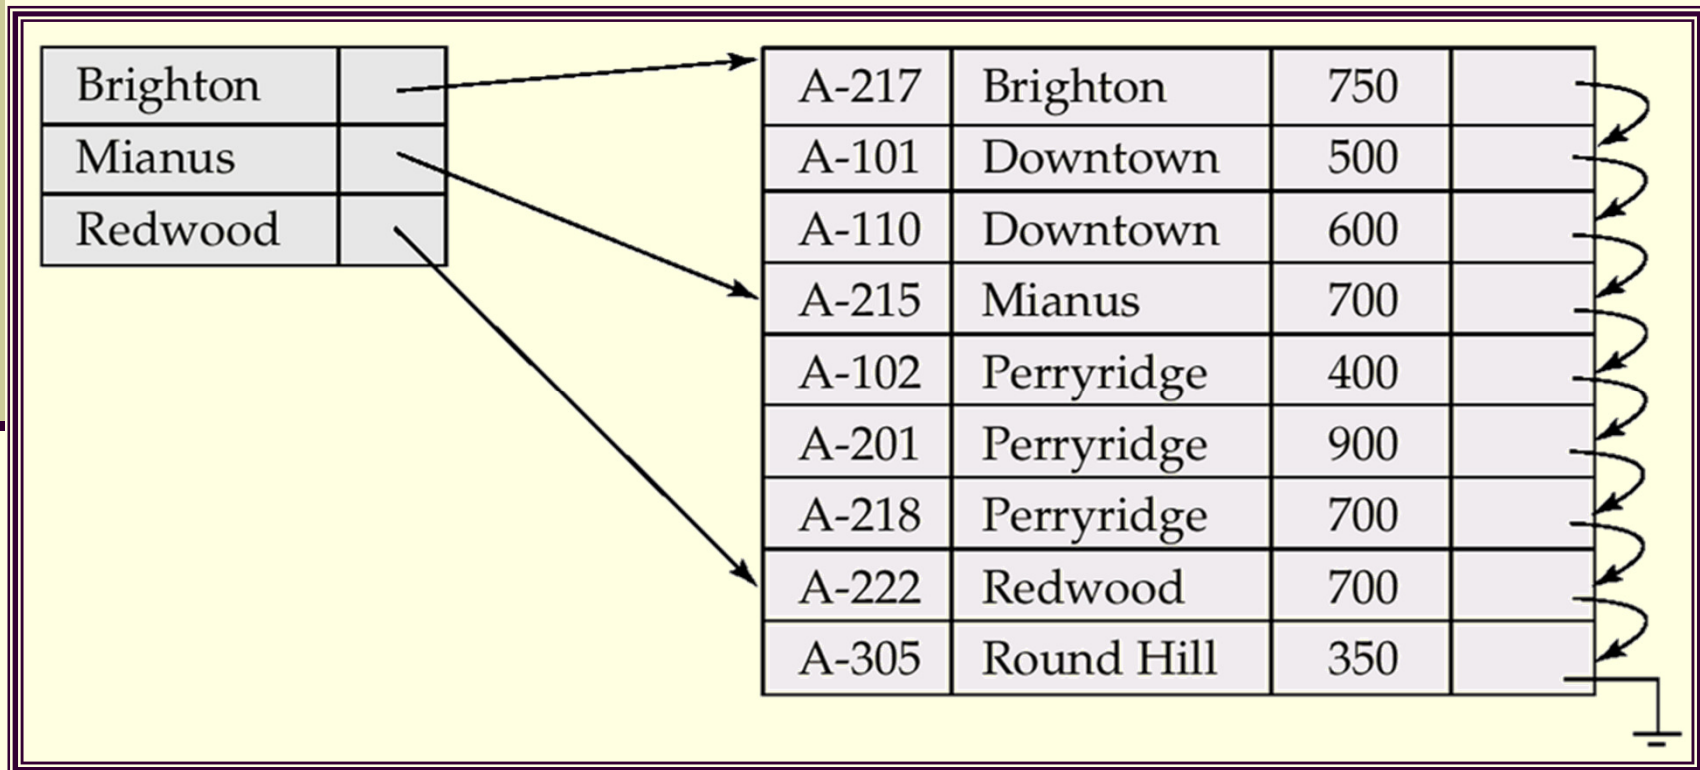
Dense Index Files

- **Dense index** — Index record appears for every search-key value in the file.



Sparse Index Files

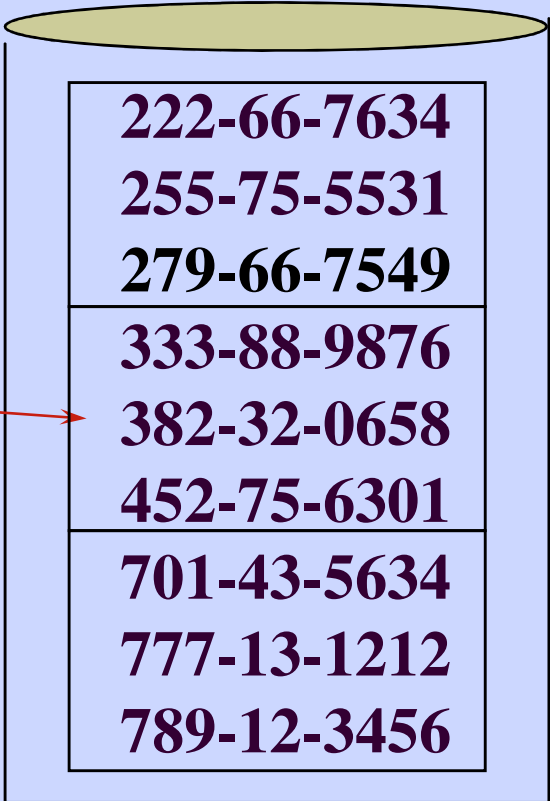
- **Sparse index** — Index record appears for only some search-key value in the file.



Indexed Sequential Access: Fast

Find record with key **382-32-0658**

<u>Key</u>	<u>Cyl. Trck Sect.</u>		
222-66-7634	3	10	2
333-88-9876	3	10	3
701-43-5634	3	10	4



222-66-7634
255-75-5531
279-66-7549
333-88-9876
382-32-0658
452-75-6301
701-43-5634
777-13-1212
789-12-3456


$333-88-9876 < \mathbf{382-32-0658} > 701-43-5634$

Search Cyl. 3, Trck 10, Sect. 3 sequentially.

Multilevel Index

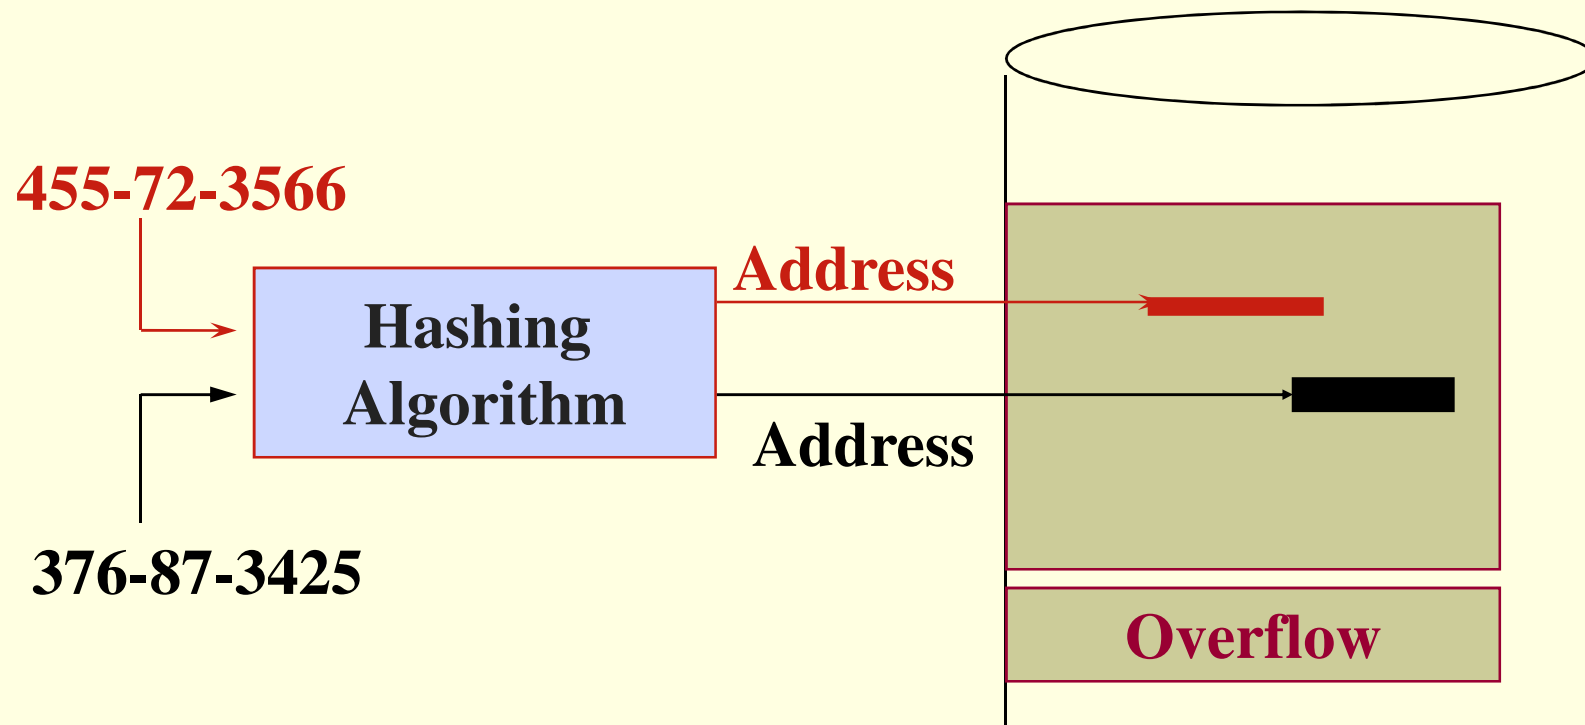
- If primary index does not fit in memory, access becomes expensive.
- To reduce number of disk accesses to index records, treat primary index kept on disk as a sequential file and construct a sparse index on it.
 - outer index – a sparse index of primary index
 - inner index – the primary index file
- If even outer index is too large to fit in main memory, yet another level of index can be created, and so on.

Multilevel Index (Cont.)



Direct or Hashed Access

- A portion of disk space is reserved
- A “hashing” algorithm computes record address



Static Hashing

- A **bucket** is a unit of storage containing one or more records (a bucket is typically a disk block).
- In a **hash file organization** we obtain the bucket of a record directly from its search-key value using a **hash function**.
- Hash function h is a function from the set of all search-key values K to the set of all bucket addresses B .
- Hash function is used to locate records for access, insertion as well as deletion.
- Records with different search-key values may be mapped to the same bucket; thus entire bucket has to be searched sequentially to locate a record.

E.g. Hash File Organization

A 1
B 2
C 3
D 4
E 5
G 7
H 8
I 9
L 12
N 14
O 15
P 16
R 18
T 20
U 21
Y 25

- There are 10 buckets,
- The binary representation of the i th character is assumed to be the integer i .
- The hash function returns the sum of the binary representations of the characters modulo 10

$$\begin{aligned} \text{e.g. } h(\text{Perryridge}) &= \text{MOD}(16+5+18+18+25+18+9+4+7+5) \\ &= \text{MOD}(125) = 5 \end{aligned}$$

$$\begin{aligned} h(\text{Round Hill}) &= \text{MOD}(18+15+21+14+4+8+9+12+12) \\ &= \text{MOD}(113) = 3 \end{aligned}$$

$$\begin{aligned} h(\text{Brighton}) &= \text{MOD}(2+18+9+7+8+20+15+14) \\ &= \text{MOD}(93) = 3 \end{aligned}$$

Example of Hash File Organization

**Hash file organization
of *account* file, using
branch-name as key**

bucket 0			
bucket 1			
bucket 2			
bucket 3	A-217	Brighton	750
	A-305	Round Hill	350
bucket 4	A-222	Redwood	700
bucket 5	A-102	Perryridge	400
	A-201	Perryridge	900
	A-218	Perryridge	700
bucket 6			
bucket 7	A-215	Mianus	700
bucket 8	A-101	Downtown	500
	A-110	Downtown	600
bucket 9			

Hash Functions


- Worst has function maps all search-key values to the same bucket; this makes access time proportional to the number of search-key values in the file.
- An ideal hash function is **uniform**, i.e., each bucket is assigned the same number of search-key values from the set of *all* possible values.
- Ideal hash function is **random**, so each bucket will have the same number of records assigned to it irrespective of the *actual distribution* of search-key values in the file.
- Typical hash functions perform computation on the internal binary representation of the search-key.

Handling of Bucket Overflows

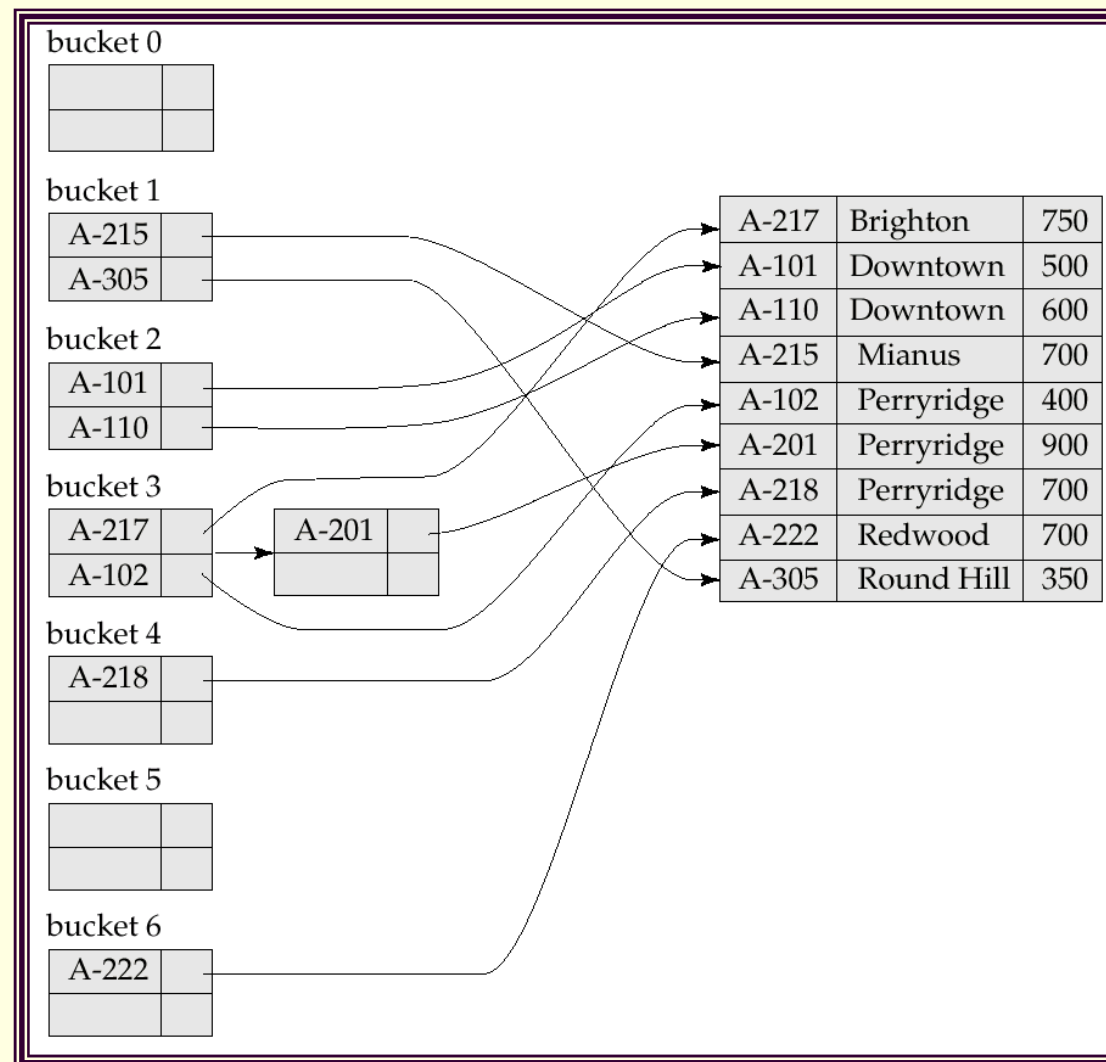
- Bucket overflow can occur because of
 - Insufficient buckets
 - Skew in distribution of records. This can occur due to two reasons:
 - multiple records have same search-key value
 - chosen hash function produces non-uniform distribution of key values
- Although the probability of bucket overflow can be reduced, it cannot be eliminated; it is handled by using *overflow buckets*.

Handling of Bucket Overflows (Cont.)

- **Overflow chaining** – the overflow buckets of a given bucket are chained together in a linked list.



E.g. of Secondary Hash Index



Deficiencies of Static Hashing

- In static hashing, function h maps search-key values to a fixed set of B of bucket addresses.
 - Databases grow with time. If initial number of buckets is too small, performance will degrade due to too much overflows.
 - If file size at some point in the future is anticipated and number of buckets allocated accordingly, significant amount of space will be wasted initially.
 - If database shrinks, again space will be wasted.
 - One option is periodic re-organization of the file with a new hash function, but it is very expensive.

Dynamic Hashing

- Good for database that grows and shrinks in size
- Allows the hash function to be modified dynamically
- **Extendable hashing** – one form of dynamic hashing


Secondary Index

- Provide access via non-key attributes
- Three data structures discussed here:
 - Linked lists: embedded pointers
 - Inverted lists: cross-index tables
 - B-Trees

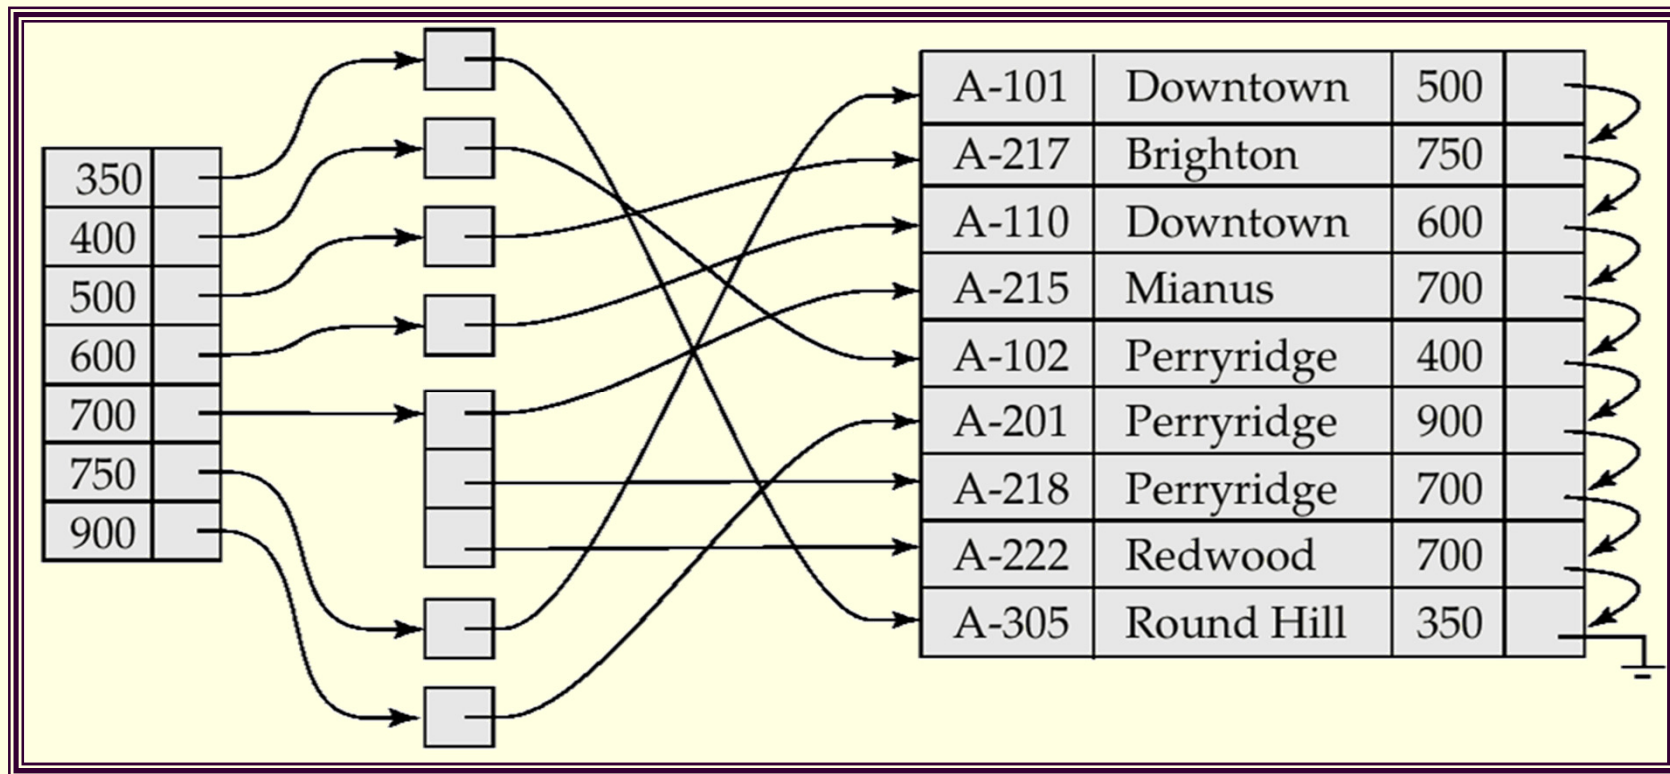
Index

- Primary mechanism to get improved performance on a database
- Persistent data structure, stored in database
- Many interesting implementation issues
 - But we are focusing on user/application perspective

Index



Secondary Index on *balance* field of *account*



Primary and Secondary Indices

- Secondary indices have to be dense.
- Indices offer substantial benefits when searching for records.
- When a file is modified, every index on the file must be updated, Updating indices imposes overhead on database modification.
- Sequential scan using primary index is efficient, but a sequential scan using a secondary index is expensive
 - each record access may fetch a new block from disk

Index

- Difference between full table scans and immediate location of tuples
 - Orders of magnitude performance difference
- Underline Data Structures
 - Balance Trees (B Trees, B+ Trees)
 - Hash Tables

Index – Key Field

```
Select sName  
From Student  
Where sID = 18942
```

- Many DBMS's build index automatically on **PRIMARY KEY** (and sometimes **UNIQUE**) attributes **sID**

Index – Non-key Fields

```
Select sID  
From Student  
Where sName = 'Mary' And GPA > 3.9
```

- Index on sName
- Index on GPA
- Index on both sName and GPA

Index – Join Tables

```
Select sName, cName  
From Student, Apply  
Where Student.sID = Apply.sID
```


Question

- Consider the following query:

```
Select * from Apply, College
Where Apply.cName = College.cName and
       Apply.major = 'CS' and
       College.enrollment < 5000)
```

Which of the following indices could **NOT** be useful in speeding up query execution?

- Tree-based index on **Apply.cName**
- Hash-based index on **College.cName**
- Hash-based index on **Apply.major**
- Hash-based index on **College.enrollment**

Question

- Consider the following query:

```
Select * from Student, Apply, College
Where Student.sID = Apply.sID and
      Apply.cName = College.cName and
      Student.GPA > 1.5 and
      College.cName < 'Cornell'
```

Suppose we are allowed to create two indexes, and assume all indexes are tree-based. Which two indexes would be most useful for speeding up query execution?

- Apply.sID, Student.GPA
- Apply.cName, College.cName
- Student.sID, Student.GPA
- Student.sID, College.cName

Downsides of Index

- Extra Space
 - Marginal
- Index Creation
 - Medium
- Index Maintenance
 - Can often set off benefits

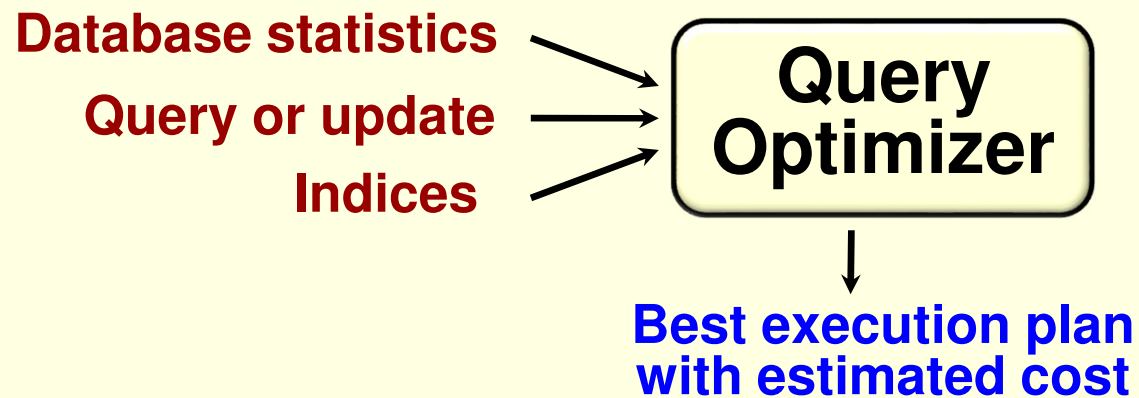
When to Create an Index?

Benefit of an index depends on:

- Size of table (and possibly layout)
- Data distributions
- Query vs. update load

Physical design advisors

- Input
 - Database (statistics) and workload
- Output
 - Recommended indices



SQL - Create & Drop Index

- Create Index **IndexName** on **T (A)**
- Create Index **IndexName** on **T (A1, A2, ..., An)**
- Create Unique Index **IndexName** on **T (A)**
- Drop Index **IndexName**

Indices


- SQL allows you to index database tables, making it possible to quickly seek to records without performing a full table scan first and thus significantly speeding up query execution.
- You can have many indexes (16 in MySQL) per table, and SQL also supports multi-column indices and full-text search indices.

Selecting Records, e.g.

- **Primary index** (records ordered on a key field)


$T1 \leftarrow \sigma_{\text{empno}='12345'}(\text{Employee})$ (single record)

$T1 \leftarrow \sigma_{\text{dno} \geq '5'}(\text{Department})$ (multiple records)




- **Clustering Index** (records ordered on a non key field)

$T1 \leftarrow \sigma_{\text{dno}='5'}(\text{Employee})$ (multiple records)



- **B⁺-Tree Index** (secondary index on equality comparison)


$T1 \leftarrow \sigma_{\text{salary} \geq 30000 \text{ and } \text{salary} \leq 35000}(\text{Employee})$ (multiple records)



Cost of Operations

	(a) Scan	(b) Equality	(c) Range	(d) Insert	(e) Delete
(1) Heap	BD	$0.5BD$	BD	$2D$	Search + D
(2) Sorted	BD	$D \log_2 B$	$D \log_2 B + \# \text{ matches}$	Search + BD	Search + BD
(3) Clustered	$1.5BD$	$D \log_F 1.5B$	$D \log_2 1.5B + \# \text{ matches}$	Search + D	Search + D
(4) Unclustered Tree index	$BD(R+0.15)$	$D(1 + \log_F 0.15B)$	$D \log_F 0.15B + \# \text{ matches}$	$D(3 + \log_F 0.15B)$	Search + $2D$
(5) Unclustered Hash index	$BD(R+0.125)$	$2D$	BD	$4D$	Search + $2D$

B: The number of data pages
R: Number of records per page
D: (Average) time to read or write disk page

 Several assumptions underlie these (rough) estimates!